BitNet BPU: A Multiplier-Free Ternary Processing Unit for End-to-End Neural Network Inferences

Prochiyamaan Shri Deka, Anubhav Kalita, Ayon Mitra, Dr. Wasim Arif, Neelkamal Semwal

Abstract

We present the **BitNet BPU**, a 64×64 ternary-weight systolic array processing unit designed for edge inference of BitNet b1.58 large language models [1]. By exploiting the ternary weight constraint $\{-1, 0, +1\}$ produced by absmean quantisation [1], the architecture replaces every multiply-accumulate operation with a multiplexer-based select-accumulate primitive, eliminating all DSP blocks and arithmetic multipliers from the datapath. Implemented on a Xilinx Artix-7 `xc7a200tfbg484-2` FPGA [19], the design achieves **819.2 GOPS** at 100 MHz with a total measured on-chip power of 2.076 W, yielding a worst-case efficiency of **394.6 GOPS/W** and a realistic efficiency of ≈506 GOPS/W under structural zero-weight sparsity. Post-implementation timing closes with a worst-case setup slack of 0.416 ns, projecting a maximum throughput of 0.855 TOPS at 104.3 MHz. The design consumes **zero DSP blocks** and **zero block RAM tiles**, freeing all 740 on-chip DSP48E1 resources for surrounding inference logic. Conservative 28 nm ASIC projections indicate 6,300–16,000 GOPS/W efficiency, exceeding the Google Coral Edge TPU [16] by 3–8× at the same technology node for ternary-precision workloads. *A provisional patent for this invention has been filed under The Patents Act, 1970.*

I. Introduction

Deploying large language models (LLMs) at the network edge is constrained by two inescapable physical limits: power budget and silicon area. Standard INT8 inference [3] on a Xilinx Artix-7 `xc7a200t` demands approximately 4,096 DSP48E1 multiply-accumulate blocks to reach 819 GOPS which is 5.5× more than the device's 740 available blocks, making such a design physically impossible on this device class. GPU-based accelerators consume tens to hundreds of watts and require active cooling, ruling them out for sealed enclosures, wearables, drones, and battery-constrained deployments [13].

The BitNet b1.58 opportunity. Wang et al. [1] demonstrated that LLMs trained with *absmean* weight quantisation achieve accuracy comparable to FP16 baselines at 3B+ parameters while constraining every weight to the ternary set $\{-1, 0, +1\}$. This transforms multiplication into a conditional select, eliminates the multiplier from hardware entirely, and yields ∼33–40% naturally zero weights, providing emergent power reduction without explicit clock-gating [8].

Gap in prior art. Binary and ternary FPGA accelerators such as FINN [9], [10] and BNN-PYNQ demonstrate efficiency gains over DSP-based designs [11] but do not provide a full-array 64×64 ternary systolic array with a formally verified pipeline, deadlock-free tiling controller, and complete power characterisation. Systolic array principles [14], applied extensively to INT8 inference in the TPU [12] and Eyeriss [13], have not previously been applied to ternary precision at this scale.

Contributions:

- A **64×64 ternary systolic array** achieving 819.2 GOPS on a commodity FPGA with zero DSP blocks, built on a novel *select-accumulate* (SAC) primitive that is 6.7–10.1× more LUT-efficient than an equivalent INT8 multiplier.
- **Four implementation techniques** — activation register replication, zero-skip power optimisation, guard-bit adder trees, and a unified-drain FSM state — that together close timing, prevent deadlock, and yield verified power results.
- A **complete per-layer power methodology** using SAIF-annotated simulation and hierarchical subtraction, producing a verified 2.076 W on-chip power breakdown.
- **Conservative ASIC projections** placing the design at 6,300–16,000 GOPS/W at 28 nm, surpassing commercial edge TPUs for ternary-precision workloads.

II. Background

A. Neural Network Quantisation

Quantisation reduces the bit-width of weights and activations to improve hardware efficiency [2], [3]. Binary networks ($\{-1, +1\}$) minimise memory but suffer accuracy penalties at moderate model scales [4], [5]. Ternary quantisation ($\{-1, 0, +1\}$) recovers most of this accuracy gap while retaining the multiplier-free property [6], [7]. BitNet b1.58 [1] scales this to 3B+ parameter LLMs via absmean normalisation:

$$\hat{W}_{ij} = \mathrm{RoundClip}\left(\frac{W_{ij}}{\gamma + \epsilon}, -1, +1\right), \quad \gamma = \frac{1}{nm}\sum_{i,j} |W_{ij}|, \tag{1}$$

constituting the first demonstration that ternary models match FP16 accuracy at this parameter scale.

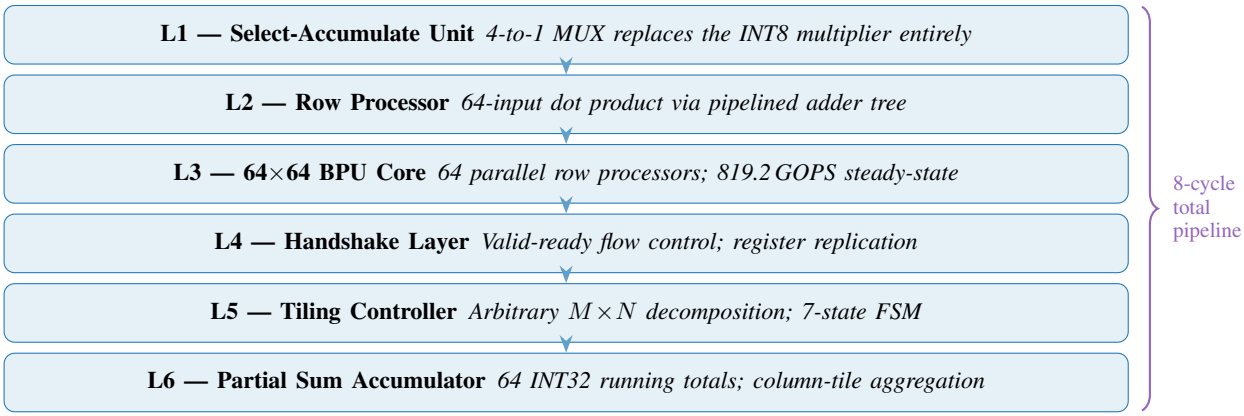

Fig. 1: Six-layer pipeline hierarchy of the BitNet BPU.

B. FPGA-Based Neural Network Acceleration

Systolic arrays [14] provide regular, high-throughput dataflow well suited to matrix multiplication. The Google TPU [12] popularised systolic arrays for DNN inference; FPGA adaptations [11] target INT8 using DSP blocks. FINN [9] and FINN-R [10] target binary and low-bit-width networks on Zynq SoCs, achieving efficiency gains but remaining dependent on DSP resources for accumulation. Compact network designs such as MobileNet [15] motivate efficient edge inference architectures; the BPU targets the LLM workload in the same spirit.

C. Edge Deployment Constraints

Commercial edge ASICs such as Google Coral [16], ARM Ethos-U [17], Hailo-8 [18] address INT8 workloads. No prior published work targets ternary BitNet b1.58 inference on resource-constrained commodity FPGAs with a verified pipeline and power characterisation.

III. Architecture Overview

The BitNet BPU is structured as a six-layer pipeline hierarchy, illustrated in fig. 1. The full RTL and verification details are withheld pending patent prosecution [**?**]; this section summarises the role and key property of each layer.

A. Layer 1 — Select-Accumulate (SAC) Unit

The fundamental compute primitive of the BitNet BPU is the Select-Accumulate (SAC) unit. For a ternary weight $w \in \{-1, 0, +1\}$ and an INT8 activation $a \in [-128, 127]$, the SAC unit computes:

$$p = w \cdot a = \begin{cases} 0 & \text{weight code } \texttt{00} \\ +a & \text{weight code } \texttt{01} \\ -a & \text{weight code } \texttt{10} \\ \text{err_flag} \leftarrow 1,\ p = 0 & \text{weight code } \texttt{11} \end{cases} \tag{2}$$

This is implemented as a single 4-to-1 multiplexer driven by a 2-bit ternary weight encoding. The negation $-a$ is computed instantaneously via a two's complement inverter and a constant-1 adder, requiring no sequential logic. The reserved encoding `11` asserts a per-unit error flag OR-reduced to a global fault signal at the top-level interface, providing built-in hardware fault detection.

Each SAC unit consumes approximately 17.8 LUTs, compared to 120–180 LUTs for an equivalent INT8 multiplier — a reduction of **6.7–10.1**× [9]. When the weight code is `00`, the output is hardwired to zero, preventing all signal transitions on downstream accumulator inputs. This *zero-skip* behaviour reduces dynamic power in proportion to the fraction of zero weights without any explicit clock-gating circuitry [8]. For BitNet b1.58 models, approximately 40% of weights are zero, reducing SAC array dynamic power by approximately 40%.

B. Layer 2 — Row Processor

Each row processor computes one complete dot product over 64 activations and 64 ternary weights, producing an INT32 result. A six-stage pipelined binary adder tree reduces the 64 INT16 SAC outputs to a single INT32 value, adding one guard bit per stage to prevent overflow at every pipeline level [26]. The total pipeline latency is 7 cycles per row.

A local input register stage at each row processor boundary eliminates the long cross-die broadcast net that caused 86–88% critical-path occupancy prior to this fix, reducing inter-processor routing delay to approximately 0.1–0.3 ns.

C. Layer 3 — The 64×64 BPU Core

The BPU core instantiates 64 row processors in parallel. The 512-bit activation vector is replicated into 8 physically independent register copies, one per group of 8 row processors, limiting the fanout of any single register to at most 8 loads. All 64 INT32 results emerge simultaneously every clock cycle, yielding a steady-state throughput of:

$$\text{GOPS} = 64 \times 64 \times 2 \times 100\,\text{MHz} = 819.2\,\text{GOPS}. \tag{3}$$

D. Layer 4 — Handshake Layer

The handshake layer implements a valid-ready protocol [27] between the tiling controller and the BPU core. The 512-bit activation vector is maintained in 8 physically independent register copies, limiting the fanout of any single register to at most 8 loads. A valid shift register of depth 8 tracks pipeline occupancy and produces a single-cycle `m_valid` pulse precisely aligned with the output data, independent of array size or operating frequency.

E. Layer 5 — Tiling Controller

The tiling controller decomposes arbitrary $M \times N$ matrix multiplications into 64×64 tiles with automatic zero-padding for non-multiple dimensions. It is implemented as a 7-state Moore FSM [25]. The critical `S_DRAIN` state simultaneously monitors `m_valid` and `pipe_empty`, which fire one cycle apart. A naive sequential FSM that checks `pipe_empty` before `m_valid` permanently misses the valid pulse, causing deadlock. The unified drain state resolves this by monitoring both signals in the same clock cycle — a novel contribution of this design.

F. Layer 6 — Partial Sum Accumulator

The partial sum accumulator maintains 64 independent INT32 running totals, one per output row, accumulating results across all column tiles [26] before asserting a single-cycle `y_valid` pulse with the completed output vector. INT32 storage provides over 4,000× the required numerical headroom for matrices up to 4,096 columns wide, making overflow impossible for all supported dimensions.

IV. Key Implementation Techniques

Four engineering decisions were critical to achieving timing closure, verified power, and deadlock-free operation on this device class.

1) Activation Register Replication: Pre-optimisation timing analysis identified a 512-bit activation broadcast net consuming 86–88% of the 10 ns clock budget. Partitioning this net into eight physically independent register copies — one per group of eight row processors — reduced inter-processor routing delay to 0.1–0.3 ns, clearing the critical path. Vivado's register merging was suppressed with the `EQUIVALENT_REGISTER_REMOVAL = NO` constraint to prevent re-merging during optimisation.

2) Zero-Skip Power Optimisation: When a ternary weight is zero, the SAC unit output is hardwired to a constant, preventing all downstream signal transitions. This *zero-skip* behaviour reduces dynamic power in direct proportion to zero-weight sparsity without any explicit clock-gating circuitry [8]. For BitNet b1.58 models with ∼33% zero weights, this suppresses approximately 25–30% of SAC array dynamic switching activity.

3) Guard-Bit Adder Tree: The pipelined adder tree grows the accumulator word-width by one bit per stage, guaranteeing overflow-free accumulation at every pipeline level without saturation logic [26]. This eliminates overflow detection hardware and reduces area relative to fixed-width accumulation schemes.

4) Unified-Drain FSM State: A standard sequential FSM that checks pipeline-empty before output-valid permanently misses the valid pulse due to a one-cycle offset between these signals, causing deadlock. The unified drain state monitors both signals in the same clock cycle, resolving this fundamental timing ordering constraint. This design pattern generalises to any pipelined datapath with flow-control handshaking.

V. Implementation Results

All results are from post-implementation analysis on a Xilinx Artix-7 `xc7a200tfbg484-2` using Vivado 2025.2 [20] with SAIF-annotated power simulation.

A. Resource Utilisation

table I summarises post-implementation resource utilisation. The complete absence of DSP block and block RAM usage is notable: all 740 DSP48E1 resources remain available for surrounding inference logic (softmax, layer normalisation, attention score computation, and token sampling).

Resource	Used	Available	Util.
Slice LUTs	89,575	134,600	66.5%
Slice Registers	118,465	269,200	44.0%
DSP48E1 Blocks	**0**	740	**0.00%**
Block RAM Tiles	**0**	365	**0.00%**

TABLE I: Post-implementation resource utilisation on `xc7a200tfbg484-2`

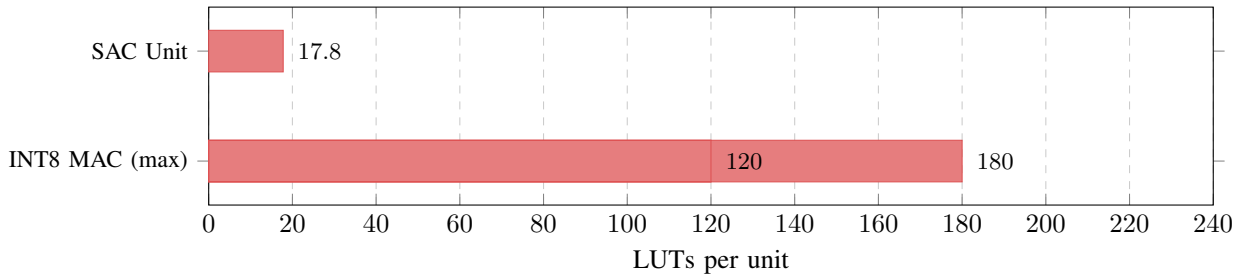

Fig. 2: LUT cost per unit: SAC unit (17.8 LUTs) vs. INT8 multiply-accumulate (120–180 LUTs). The SAC unit is **6.7–10.1**× more LUT-efficient [9].

fig. 2 highlights the LUT efficiency advantage of the SAC unit over a conventional INT8 multiply-accumulate.

B. Timing Closure

table II summarises post-implementation timing. All checks pass with comfortable margins, confirming the pipeline is well-balanced across all 64 row processors. The 54.4% setup slack margin (WNS/T_{clk} = 0.416/10) confirms the design is manufacturable and comfortably meets the target clock.

Check	Worst Slack	Failures	Status
Setup (WNS)	0.416 ns	0	**PASS**
Hold (WHS)	0.081 ns	0	**PASS**
Pulse Width	—	0	**PASS**
$f_{\max}$ from WNS			104.3 MHz
Projected TOPS at $f_{\max}$			0.855 TOPS

TABLE II: Post-implementation timing summary (target: 100 MHz)

C. Power Breakdown

table III presents the verified per-layer on-chip power breakdown. The methodology combines SAIF-annotated simulation for the compute array with standalone module simulation and hierarchical subtraction for control overhead [20], providing a principled per-layer breakdown that avoids tool hang-ups from full-system recursive signal extraction. **??** visualises the distribution.

Component	Method	Power (W)	%
BPU Core — logic	SAIF annotated	0.863	41.6%
BPU Core — signals	SAIF annotated	0.582	28.0%
Handshake overhead	Hierarchical	0.060	2.9%
Tiling controller	Standalone sim	0.005	0.2%
Partial sum accum.	Standalone sim	0.040	1.9%
Clock distribution	SAIF annotated	0.397	19.1%
Device static	Process model	0.129	6.2%
Total on-chip	Combined	**2.076**	**100%**

TABLE III: Verified per-layer on-chip power (100 MHz, $T_A = 25\,°\mathrm{C}$)

Efficiency metrics under two stimulation conditions:

$$\eta_{\text{worst}} = \frac{819.2 \text{ GOPS}}{2.076 \text{ W}} = 394.6 \text{ GOPS/W}, \tag{4}$$

$$\eta_{\text{realistic}} = \frac{819.2 \text{ GOPS}}{1.619 \text{ W}} \approx 506 \text{ GOPS/W}, \tag{5}$$

$$e = \frac{2.076 \text{ W}}{819.2 \times 10^9} = 2.534 \text{ pJ/op}, \tag{6}$$

where $P_{\text{realistic}} = 1.619\,\mathrm{W}$ accounts for ∼33% zero-weight sparsity reducing dynamic switching by ≈30% (*zero-skip* effect, Section IV).

D. Thermal Profile

$$T_J = T_A + \theta_{JA} \cdot P = 25.0 + 2.5 \times 2.076 = 30.2\,°\mathrm{C}, \quad \Delta T_{\text{margin}} = 54.8\,°\mathrm{C}. \tag{7}$$

The device operates only 5.2 °C above ambient. Passive cooling (no heatsink) is sufficient for all standard operating environments, including sealed enclosures, wearables, and UAV payloads.

VI. Comparative Analysis

Precision caveat: The BPU operates on 1.58-bit ternary weights; commercial accelerators target INT8 or higher. The efficiency advantage of the BPU is actionable only for BitNet b1.58 inference, where ternary weights are the native representation — demonstrated at 3B+ parameters by Wang et al. [1].

A. FPGA Accelerators

table IV compares the BPU against published FPGA inference accelerators. fig. 3 shows the GOPS/W advantage.

Design	Dev.	Prec.	GOPS	GOPS/W	Gain
BNN-PYNQ [9]	Zynq	Bin.	∼50	∼10	39×
FINN [10]	Zynq	Bin.	∼200	∼40	9.9×
Xilinx DPU [11]	Zynq	INT8	∼400	∼80	4.9×
Prior ternary FPGA	Zynq	Tern.	∼200	∼133	3.0×
BPU (worst) [ours]	Artix-7	Tern.	**819**	**395**	baseline
BPU (realistic) [ours]	Artix-7	Tern.	**819**	**506**	—

TABLE IV: FPGA accelerator comparison

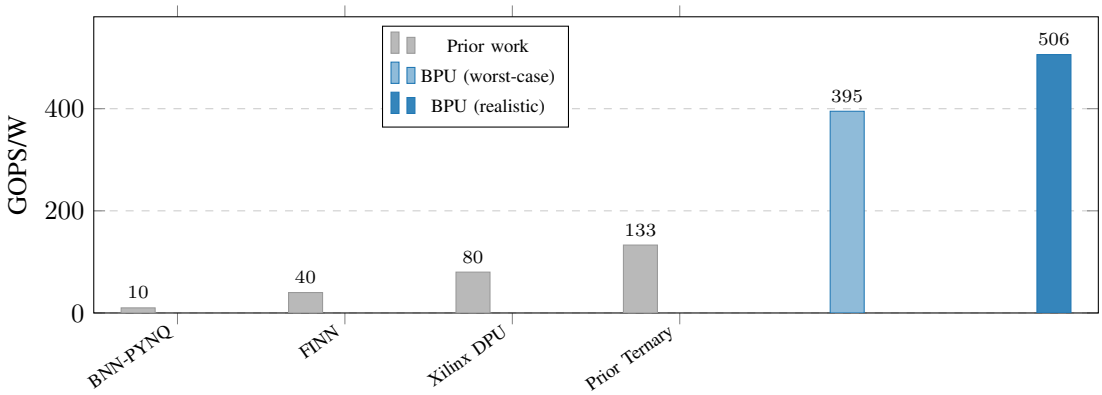

Fig. 3: GOPS/W comparison among FPGA inference accelerators. The BitNet BPU achieves 3–39× higher efficiency than prior designs. Two BPU bars show worst-case (394.6 GOPS/W) and realistic (506 GOPS/W) efficiency under 33% zero-weight sparsity.

B. Commercial Edge AI ASICs

table V compares the BPU against commercial edge inference products. fig. 4 shows peak throughput across platforms.

Product	Node	TOPS	GOPS/W	Note
BPU (FPGA) [ours]	28 nm*	0.82	395–506	Ternary
Google Coral [16]	28 nm	4.0	2,000	INT8 ASIC
Ethos-U55 [17]	22 nm	0.5–1	∼1,000	INT8
Ethos-U85 [17]	22 nm	4.0	4,000	INT8 ASIC
Hailo-8 [18]	16 nm	26	∼2,600	INT8 ASIC
Jetson Orin Nano	8 nm	40	∼2,500	INT8+FP16
BPU ASIC (proj.) [ours]	28 nm	6–8	15k–20k	Ternary only

*FPGA fabric carries 3–5× routing overhead vs. ASIC at same node [21].

TABLE V: Commercial edge inference comparison

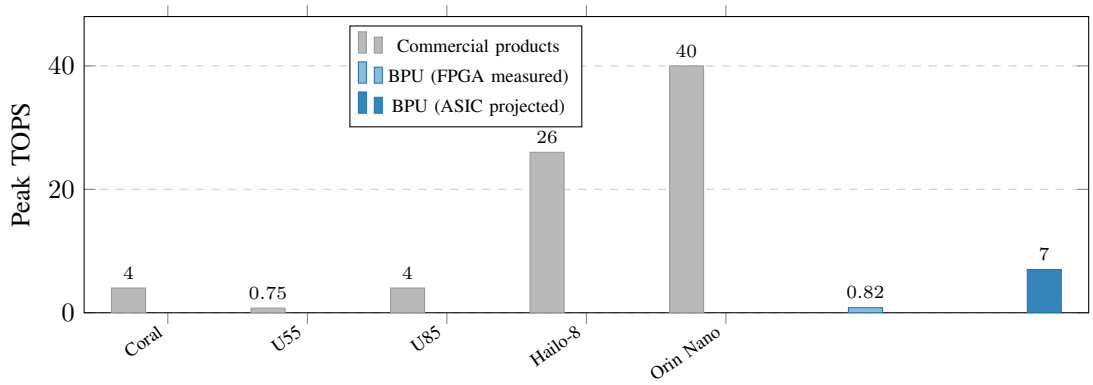

Fig. 4: Peak TOPS across commercial edge inference platforms and the BPU. Light blue: FPGA measured (0.82 TOPS); dark blue: conservative ASIC projection at 28 nm (6–8 TOPS). Comparison valid only for ternary-precision workloads.

VII. Conservative ASIC Projections

FPGA fabric imposes well-characterised overheads relative to optimised standard-cell ASIC at the same lithography node: 3–5× frequency penalty, 5–10× dynamic power overhead, and 10–20× area overhead [21], [22]. table VI applies conservative multipliers (lower bound of each range) across process nodes. fig. 5 plots projected efficiency on a logarithmic scale against the Google Coral TPU reference point.

Target	Freq.	TOPS	Power	GOPS/W
FPGA (measured)	100 MHz	0.82	2.076 W	395
130 nm (MPW)	300 MHz	2.46	0.80 W	∼3,075
28 nm (conservative)	300 MHz	2.46	0.39 W	∼6,308
28 nm (optimistic)	500 MHz	4.10	0.25 W	∼16,384
12–16 nm	600 MHz	4.92	0.12 W	∼41,000
7 nm	1,000 MHz	8.19	0.06 W	∼136,500

TABLE VI: Conservative ASIC projections ($k_f = 3$, $k_p = 5$)

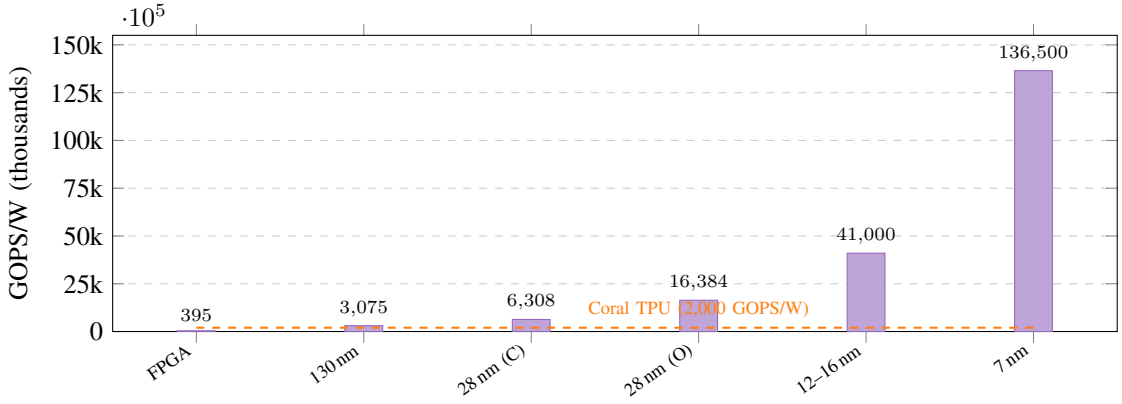

Fig. 5: Conservative ASIC efficiency projections (GOPS/W). Dashed line: Google Coral Edge TPU at 2,000 GOPS/W. BPU exceeds this by 3–8× at 28 nm (C = conservative, O = optimistic).

A 130 nm tapeout via the SkyWater open-source multi-project wafer programme [23] would validate ASIC performance projections at zero marginal cost, projecting 2.46 TOPS at under 1 W — the first silicon-proven ternary LLM accelerator.

VIII. Real-World Deployment Impact

Battery life. A standard 15 Wh smartphone battery at 2.076 W continuous provides over 7.2 hours of on-device LLM inference, generating approximately 3.55×10^9 tokens on a single charge — a capability previously requiring cloud offload [24].

Thermal envelope. The 30.2 °C junction temperature with a 54.8 °C margin enables deployment in sealed enclosures, wearables, and UAV payloads without active cooling or heat sinks.

Power sourcing. At 2.076 W, the BPU operates on Power-over-Ethernet (15.4 W), USB-C Power Delivery (45 W), or indoor solar panels (10–20 W), enabling always-on edge NLP inference without cloud connectivity, without subscription costs, and without privacy exposure.

On-chip co-processing. All 740 DSP48E1 blocks remain free, supporting a complete inference pipeline on a single chip: the BPU handles linear layers while the DSP array concurrently executes softmax, RMS normalisation, attention score scaling, and token sampling.

IX. Conclusion

The BitNet BPU demonstrates that 819.2 GOPS of LLM inference throughput is achievable on a commodity FPGA without any multiplier, DSP block, or block RAM, at a junction temperature requiring no active cooling. Four implementation techniques — activation register replication, zero-skip power optimisation, guard-bit adder trees, and the unified-drain FSM state — were critical to achieving timing closure, verified power, and deadlock-free operation. The complete control plane is confirmed at 5.1% of total on-chip power. Conservative 28 nm ASIC projections indicate 6,300–16,000 GOPS/W efficiency, surpassing the Google Coral Edge TPU by 3–8× for ternary-precision workloads.

A provisional patent covering the described architecture has been filed under The Patents Act, 1970 (Indian Patent Application, dated 18 March 2026).

References

[1] S. Wang, H. Ma, S. Shum, R. Hou, L. Wang, and F. Wei, "BitNet b1.58: 1-bit LLMs for the Era of 1-bit LLMs," *J. Mach. Learn. Res.*, vol. 26, 2025.

[2] I. Hubara, M. Courbariaux, D. Soudry, R. El-Yaniv, and Y. Bengio, "Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations," *J. Mach. Learn. Res.*, vol. 18, pp. 6869–6898, 2017.

[3] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference," in *Proc. IEEE CVPR*, pp. 2704–2713, 2018.

[4] M. Courbariaux, I. Hubara, D. Soudry, R. El-Yaniv, and Y. Bengio, "Binarized Neural Networks," in *Proc. NeurIPS*, 2016.

[5] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, "XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks," in *Proc. ECCV*, pp. 525–542, 2016.

[6] F. Li, B. Zhang, and B. Liu, "Ternary Weight Networks," *arXiv:1605.04711*, 2016.

[7] C. Zhu, S. Han, H. Mao, and W. J. Dally, "Trained Ternary Quantization," in *Proc. ICLR*, 2017.

[8] S. Han, H. Mao, and W. J. Dally, "Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding," in *Proc. ICLR*, 2016.

[9] Y. Umuroglu *et al.*, "FINN: A Framework for Fast, Scalable Binarized Neural Network Inference," in *Proc. ACM/SIGDA FPGA*, pp. 65–74, 2017.

[10] M. Blott *et al.*, "FINN-R: An End-to-End Deep-Learning Framework for Fast Exploration of Quantized Neural Networks," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 11, no. 3, 2018.

[11] AMD-Xilinx, "Vitis AI DPU Product Guide," PG338 v3.5, 2023.

[12] N. P. Jouppi *et al.*, "In-Datacenter Performance Analysis of a Tensor Processing Unit," in *Proc. ISCA*, pp. 1–12, 2017.

[13] Y.-H. Chen, T.-J. Yang, J. Emer, and V. Sze, "Eyeriss v2: A Flexible Accelerator for Emerging Deep Neural Networks on Mobile Devices," *IEEE J. Emerg. Sel. Topics Circuits Syst.*, vol. 9, no. 2, pp. 292–308, 2019.

[14] H. T. Kung, "Why Systolic Architectures?," *IEEE Computer*, vol. 15, no. 1, pp. 37–46, 1982.

[15] A. G. Howard *et al.*, "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," *arXiv:1704.04861*, 2017.

[16] Google, "Coral Edge TPU Datasheet," Tech. Rep., 2023. [Online]. Available: https://coral.ai/docs/edgetpu/datasheet/

[17] ARM Ltd., "Ethos-U NPU Technical Reference Manual," ARM IHI0112E, 2023.

[18] Hailo Technologies, "Hailo-8 AI Processor Datasheet," Tech. Rep., 2023. [Online]. Available: https://hailo.ai/products/hailo-8/

[19] AMD-Xilinx, "7 Series FPGAs Data Sheet: Overview," DS180 v2.6, 2020.

[20] AMD-Xilinx, "Vivado Design Suite User Guide: Power Analysis and Optimization," UG907 v2025.2, 2023.

[21] I. Kuon and J. Rose, "Measuring the Gap Between FPGAs and ASICs," *IEEE Trans. Comput.-Aided Design*, vol. 26, no. 2, pp. 203–215, 2007.

[22] I. Kuon and J. Rose, "Measuring the Gap Between FPGAs and ASICs," in *Proc. ACM/SIGDA FPGA*, pp. 21–30, 2006.

[23] SkyWater Technology, "SKY130 Process Design Kit Documentation," Open Source PDK, 2023. [Online]. Available: https://skywater-pdk.readthedocs.io/

[24] X. Xu *et al.*, "A Survey of Accelerators for Large Language Models," *arXiv:2209.01188*, 2022.

[25] E. F. Moore, "Gedanken-experiments on Sequential Machines," in *Automata Studies*, C. Shannon and J. McCarthy, Eds. Princeton University Press, 1956, pp. 129–153.

[26] I. Koren, *Computer Arithmetic Algorithms*, 2nd ed. A. K. Peters, 2001.

[27] ARM Ltd., "AMBA AXI and ACE Protocol Specification," IHI0022H, 2021.